



AMD CROSSFIRE DIRECTX® 11 SAMPLE

ALEX KHARLAMOV
ANAS LASRAM



TABLE OF CONTENTS



- ▲ AMD Crossfire Driver Behavior
- ▲ AMD Crossfire API
- ▲ Crossfire DirectX® 11 Sample
- ▲ Performance



AMD CROSSFIRE DRIVER

DEFAULT BEHAVIOR



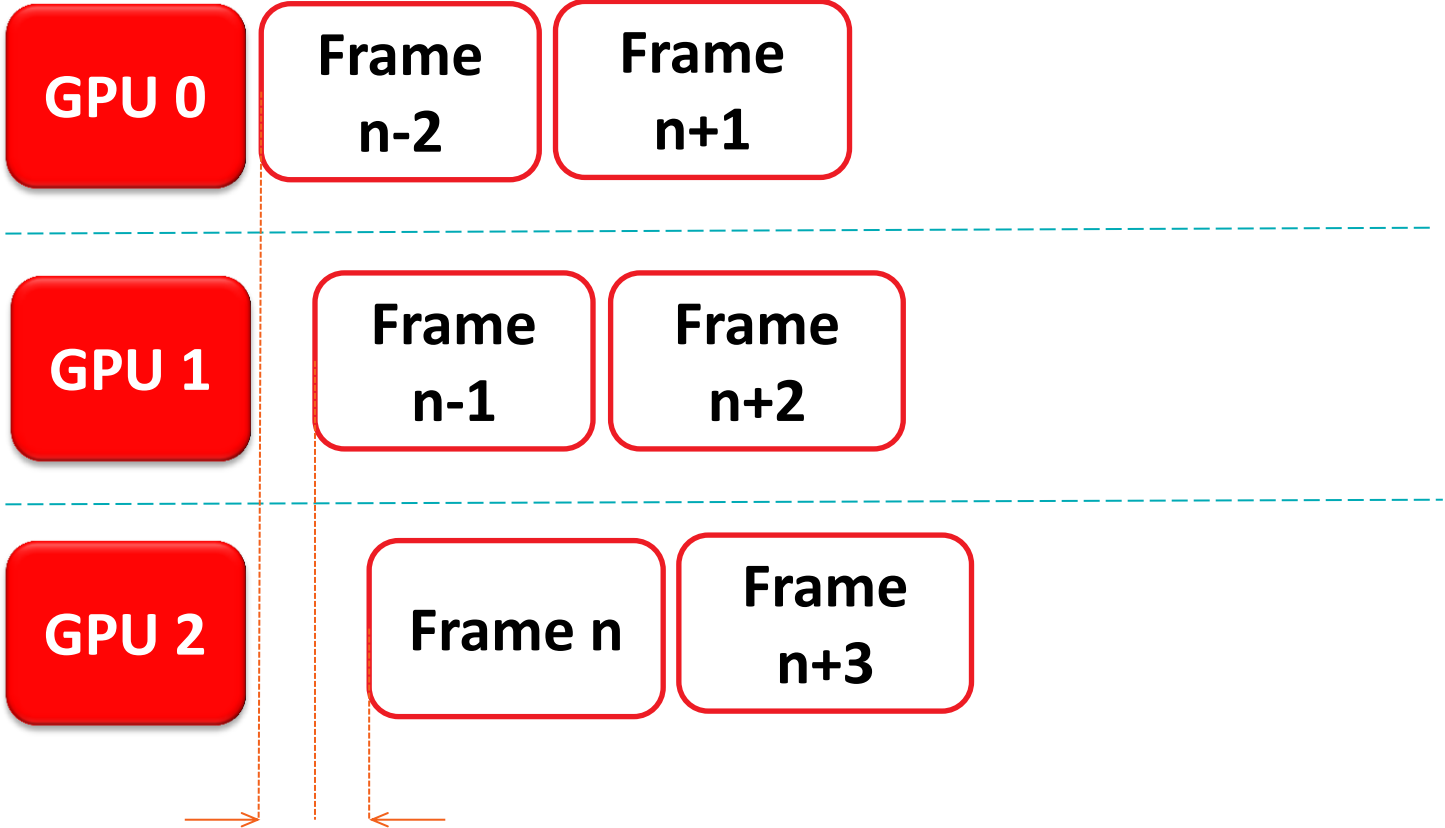
AMD CROSSFIRE



OPTIMAL APPLICATION FRAME SUBMISSION ON MULTI GPU SYSTEMS



Maximize frame rendering overlap between frames for best M-GPU scaling



Lost performance opportunity

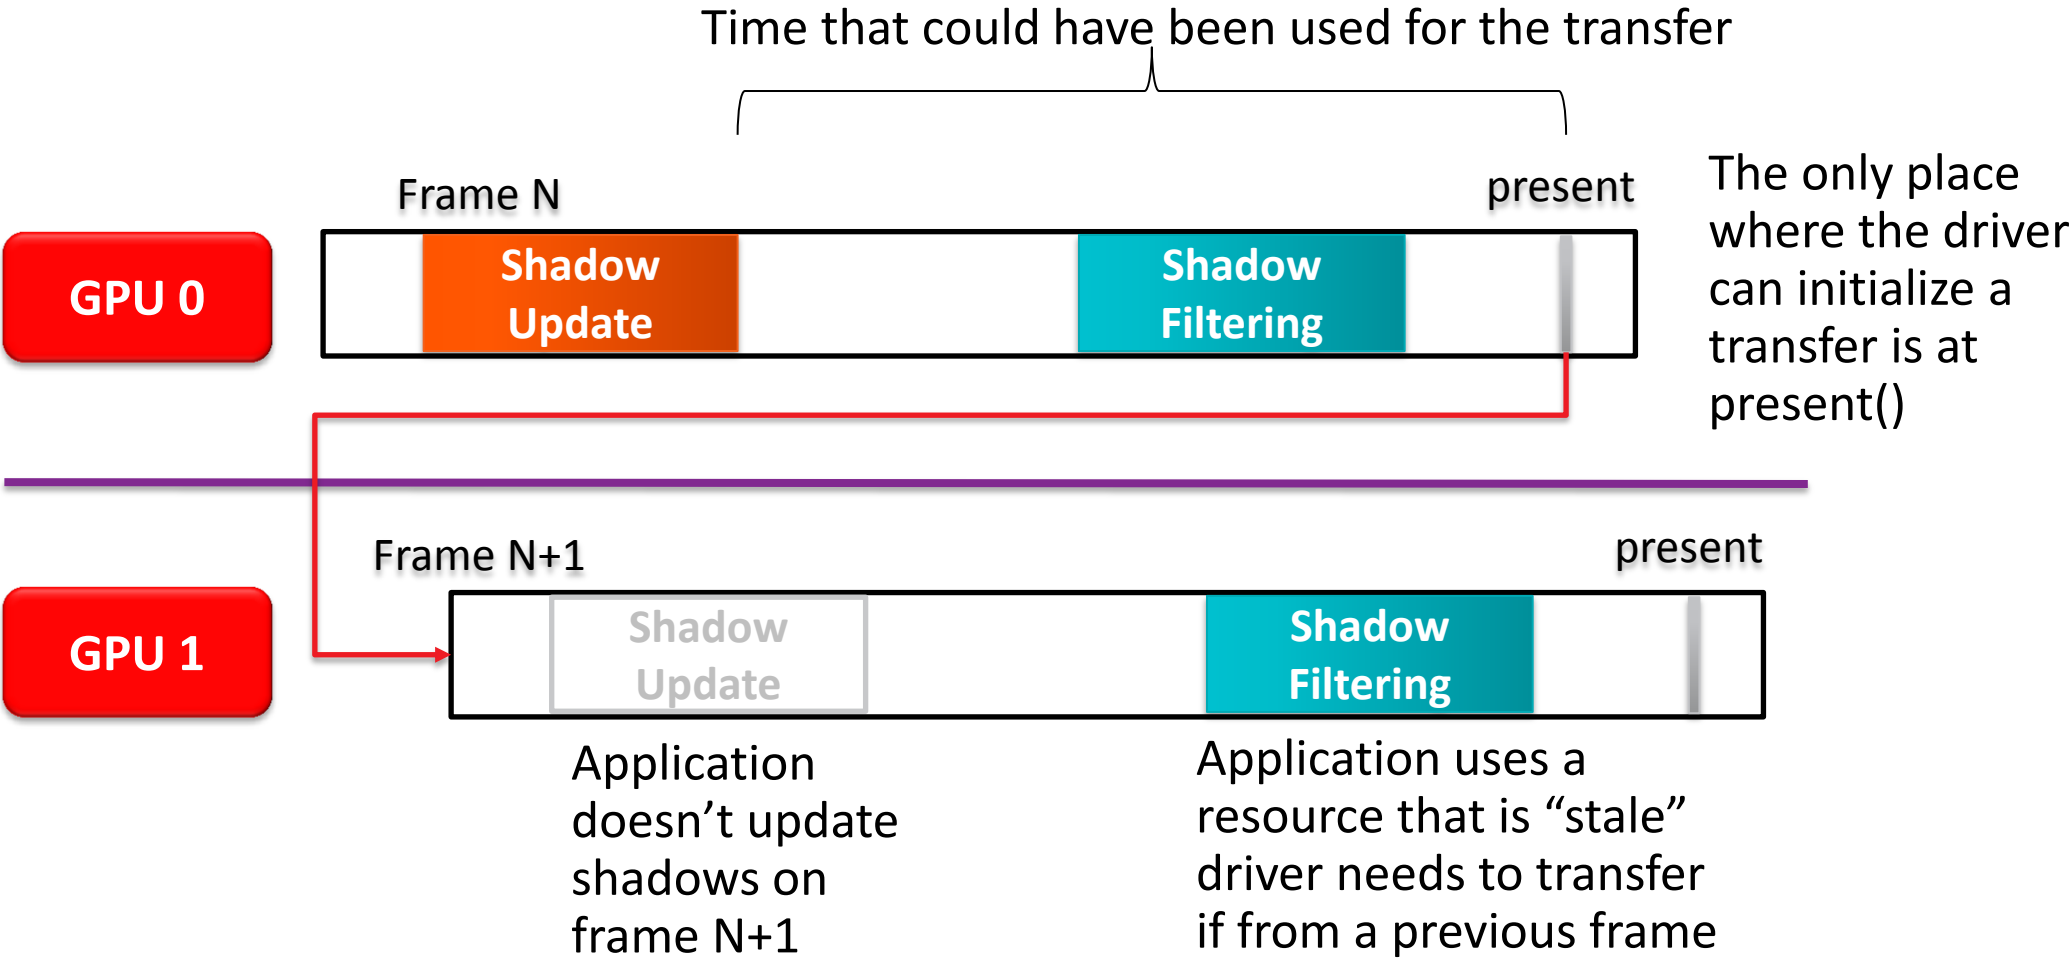
▲ **Heuristic 1:** A resource that is updated before it gets used within a frame is not stale

▲ **Heuristic 2:** Let N be the number of GPUs in the system. A resource that is updated for N frames in a row before it gets used is not stale

AMD CROSSFIRE: DEFAULT BEHAVIOR

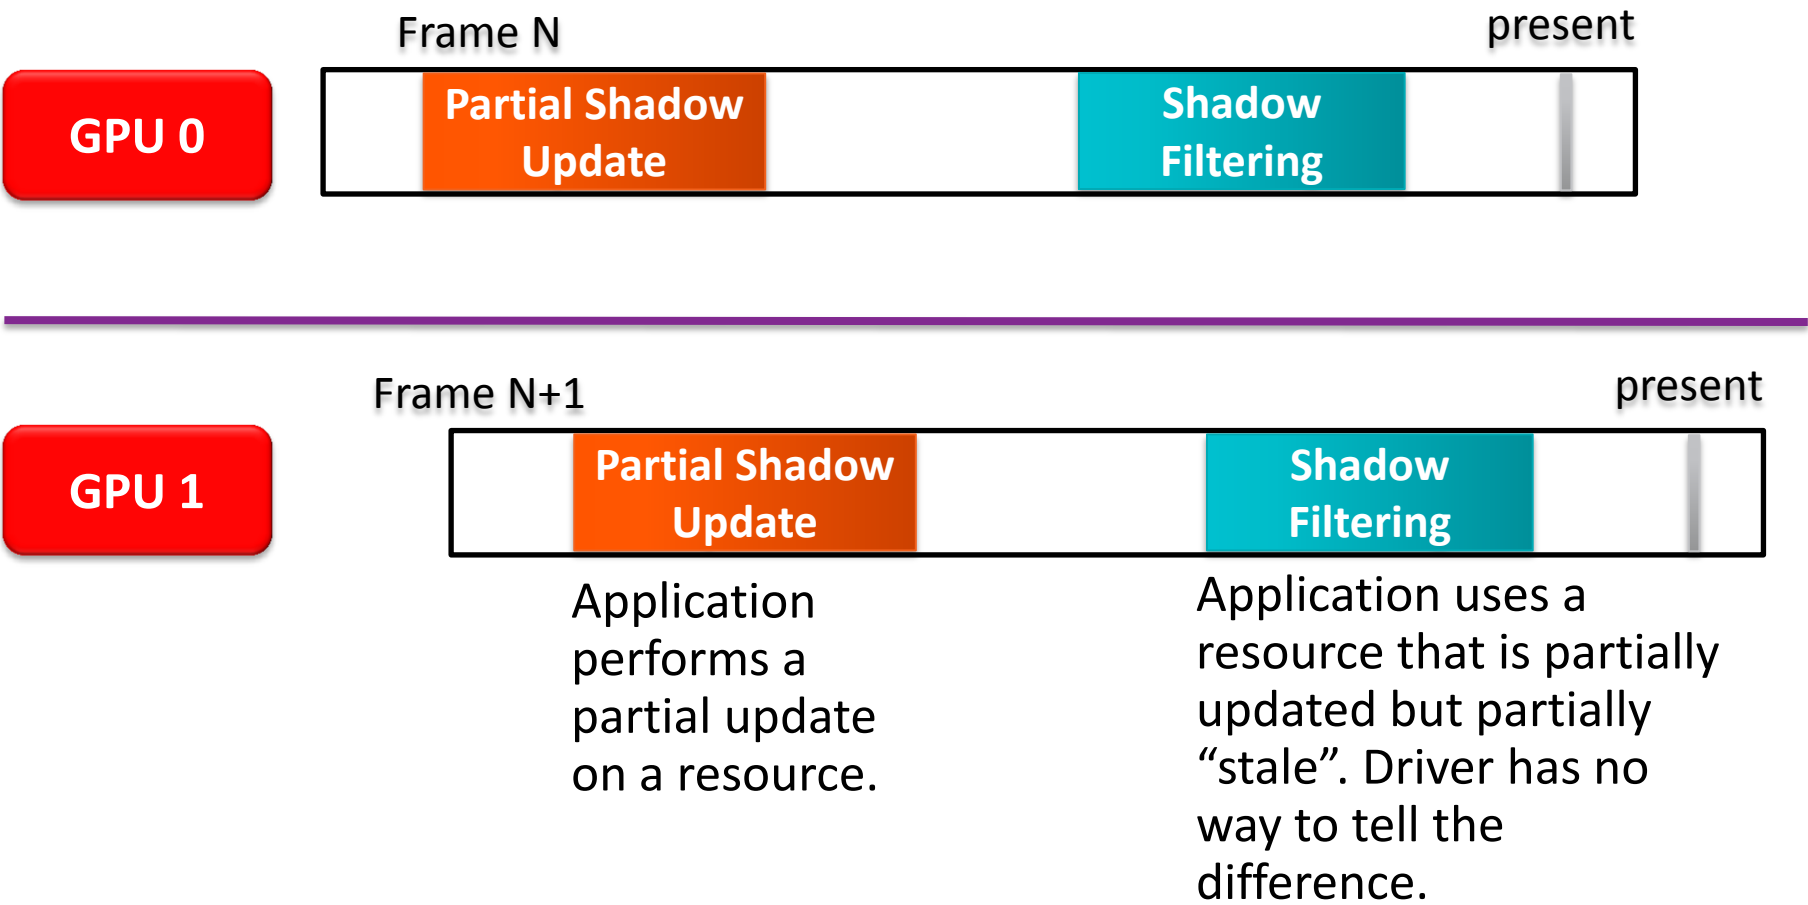


IF RESOURCE IS NOT UPDATED : TRANSFERS



AMD CROSSFIRE: DEFAULT BEHAVIOR

IF RESOURCE IS PARTIALLY UPDATED : CORRUPTION





AMD CROSSFIRE API 



▲ Driver extension headers

- Plain headers
- Application is responsible for loading driver DLLs and querying interfaces

```
class IAmDxExtAfrControl : public IAmDxExtInterface
{
public:
    // Version information
    virtual HRESULT GetExtensionVersion(AmDxExtVersion* pExtVer) = 0;

    // Enable AFR control API. Should be called before creating resources to signal intent to
    // use the AFR control API.
    virtual VOID EnableAfrControl() = 0;

    // Create resource methods with transfer type specifier
    virtual HRESULT CreateBuffer(const D3D11_BUFFER_DESC* pDesc,
                                const D3D11_SUBRESOURCE_DATA* pInitialData,
                                ID3D11Buffer** ppBuffer,
                                AmDxTransferType transferType) = 0;
    virtual HRESULT CreateTexture1D(const D3D11_TEXTURE1D_DESC* pDesc,
                                    const D3D11_SUBRESOURCE_DATA* pInitialData,
                                    ID3D11Texture1D** ppTexture1D,
                                    AmDxTransferType transferType) = 0;
    virtual HRESULT CreateTexture2D(const D3D11_TEXTURE2D_DESC* pDesc,
                                    const D3D11_SUBRESOURCE_DATA* pInitialData,
                                    ID3D11Texture2D** ppTexture2D,
                                    AmDxTransferType transferType) = 0;
    virtual HRESULT CreateTexture3D(const D3D11_TEXTURE3D_DESC* pDesc,
                                    const D3D11_SUBRESOURCE_DATA* pInitialData,
                                    ID3D11Texture3D** ppTexture3D,
                                    AmDxTransferType transferType) = 0;

    // Return the number of GPUs the UMD will use for rendering (as opposed to the number in the
    // system). Using this allows the UMD app profile to clamp the number of GPUs returned.
    virtual UINT GetNumRenderGpus() = 0;

    // This will initiate a transfer for AmDxTransferApp1StepP2P,
    // AmDxTransferApp2StepNoBroadcast, and AmDxTransferApp2StepWithBroadcast.
    // If a subregion is specified and something other than AmDxTransferApp2StepWithBroadcast is used
    // for more than 2 GPUs, it is the app's responsibility to propagate the subregions to all GPUs.
    virtual void NotifyResourceEndWrites(ID3D11Resource* pResource,
                                         const D3D11_RECT* pTransferRegion,
                                         const UINT* pSubresourceIndex) = 0;

    // This will notify the driver that the app will begin read/write access to the resource.
    virtual void NotifyResourceBeginAllAccess(ID3D11Resource* pResource) = 0;

    // This is used for AmDxTransferApp1StepP2P to notify when it is safe to initiate a transfer.
    // This call in frame N-(NumGpus-1) allows a 1 step P2P in frame N to start.
    // This should be called after NotifyResourceEndWrites.
    virtual void NotifyResourceEndAllAccess(ID3D11Resource* pResource) = 0;
};
```

AMD CROSSFIRE API

SAMPLES USE AGS LIBRARY



▲ AGS library

- ISV developed and maintained library
- Wraps multiple driver extensions

```
AMD_AGS_API AGSReturnCode agsDriverExtensions_SetCrossfireMode( AGSContext* context, AGSCrossfireMode mode );

// Description
// Functions to create a Direct3D11 resource with the specified AFR transfer type
//
// Input params
// context - Pointer to a context.
// desc - Pointer to the D3D11 resource description.
// initialData - Optional pointer to the initializing data for the resource.
// transferType - The transfer behavior. See AGSAfrTransferType for more details.
//
// Output params
// buffer/texture - Returned pointer to the resource.
//
AMD_AGS_API AGSReturnCode agsDriverExtensions_CreateBuffer( AGSContext* context, const D3D11_BUFFER_DESC* desc, const D3D11_SUBRESOURCE_DATA* initialData, ID3D11Buffer** buffer, AGSAfrTransferType transferType );
AMD_AGS_API AGSReturnCode agsDriverExtensions_CreateTexture1D( AGSContext* context, const D3D11_TEXTURE1D_DESC* desc, const D3D11_SUBRESOURCE_DATA* initialData, ID3D11Texture1D** texture1D, AGSAfrTransferType transferType );
AMD_AGS_API AGSReturnCode agsDriverExtensions_CreateTexture2D( AGSContext* context, const D3D11_TEXTURE2D_DESC* desc, const D3D11_SUBRESOURCE_DATA* initialData, ID3D11Texture2D** texture2D, AGSAfrTransferType transferType );
AMD_AGS_API AGSReturnCode agsDriverExtensions_CreateTexture3D( AGSContext* context, const D3D11_TEXTURE3D_DESC* desc, const D3D11_SUBRESOURCE_DATA* initialData, ID3D11Texture3D** texture3D, AGSAfrTransferType transferType );

// Description
// Functions to notify the driver that we have finished writing to the resource this frame.
// This will initiate a transfer for AGS_AFR_TRANSFER_1STEP_P2P,
// AGS_AFR_TRANSFER_2STEP_NO_BROADCAST, and AGS_AFR_TRANSFER_2STEP_WITH_BROADCAST.
//
// Input params
// context - Pointer to a context.
// resource - Pointer to the resource.
// transferRegions - An array of transfer regions (can be null to specify the whole area).
// subresourceArray - An array of subresource indices (can be null to specify all subresources).
// numSubresources - The number of subresources in subresourceArray OR number of transferRegions. Use 0 to specify ALL subresources and one transferRegion (which may be null if specifying the whole area).
//
AMD_AGS_API AGSReturnCode agsDriverExtensions_NotifyResourceEndWrites( AGSContext* context, ID3D11Resource* resource, const D3D11_RECT* transferRegions, const unsigned int* subresourceArray, unsigned int numSubresources );

// Description
// This will notify the driver that the app will begin read/write access to the resource.
//
// Input params
// context - Pointer to a context.
// resource - Pointer to the resource.
//
AMD_AGS_API AGSReturnCode agsDriverExtensions_NotifyResourceBeginAllAccess( AGSContext* context, ID3D11Resource* resource );

// Description
// This is used for AGS_AFR_TRANSFER_1STEP_P2P to notify when it is safe to initiate a transfer.
// This call in frame N-(NumGpus-1) allows a 1 step P2P in frame N to start.
// This should be called after agsDriverExtensions_NotifyResourceEndWrites.
//
// Input params
// context - Pointer to a context.
// resource - Pointer to the resource.
//
AMD_AGS_API AGSReturnCode agsDriverExtensions_NotifyResourceEndAllAccess( AGSContext* context, ID3D11Resource* resource );
```

AMD CROSSFIRE API



Resource Creation Flag	Description
DEFAULT	Behavior depends on Radeon Settings Crossfire mode
DISABLE	Turn off driver resource tracking
1STEP_P2P	app controlled GPU to next GPU transfer
2STEP_NO_BROADCAST	app controlled GPU to next GPU transfer using intermediate system memory
2STEP_WITH_BROADCAST	app controlled GPU to all render GPUs transfer using intermediate system memory

Radeon Settings CrossFire mode	Implicit control: The CrossFire API is not queried by the application	Explicit control: The CrossFire API is successfully queried by the application
Disabled	The application runs in single GPU mode	The CrossFire API cannot be queried. The application runs in single GPU mode
Default mode	If the application has a driver profile it will be used. If a driver profile does not exist the application will run in single GPU mode	If the application has a driver profile it will be used. If a driver profile does not exist, the CrossFire API cannot be queried and the application will run in single GPU mode
AFR friendly	The application runs in MGPU mode but resource tracking is disabled	Resource transfers are only applied when the application notifies the driver to do so. Resources created with the flag AFR_TRANSFER_DEFAULT or through the normal D3D11 API will not be tracked by the driver
Optimize 1x1	The application runs in MGPU mode and resources will be tracked by the driver and will be rendered on each GPU if their resolution is 1x1	Resource transfers are applied when the application notifies the driver to do so. Resources created with the flag AFR_TRANSFER_DEFAULT or through the normal D3D11 API will be tracked by the driver. Resources with dimensions of 1x1 will be rendered on each GPU.
AFR Compatible	The application runs in MGPU mode and resource tracking is enabled	Resource transfers are applied when the application notifies the driver to do so. Resources created with the flag AFR_TRANSFER_DEFAULT or through the normal D3D11 API will be tracked by the driver
Use AMD pre-defined profile	The mode allows using one of the existing driver profiles and applying it to the current application	The mode allows using one of the existing driver profiles and applying it to the current application

AMD CROSSFIRE API



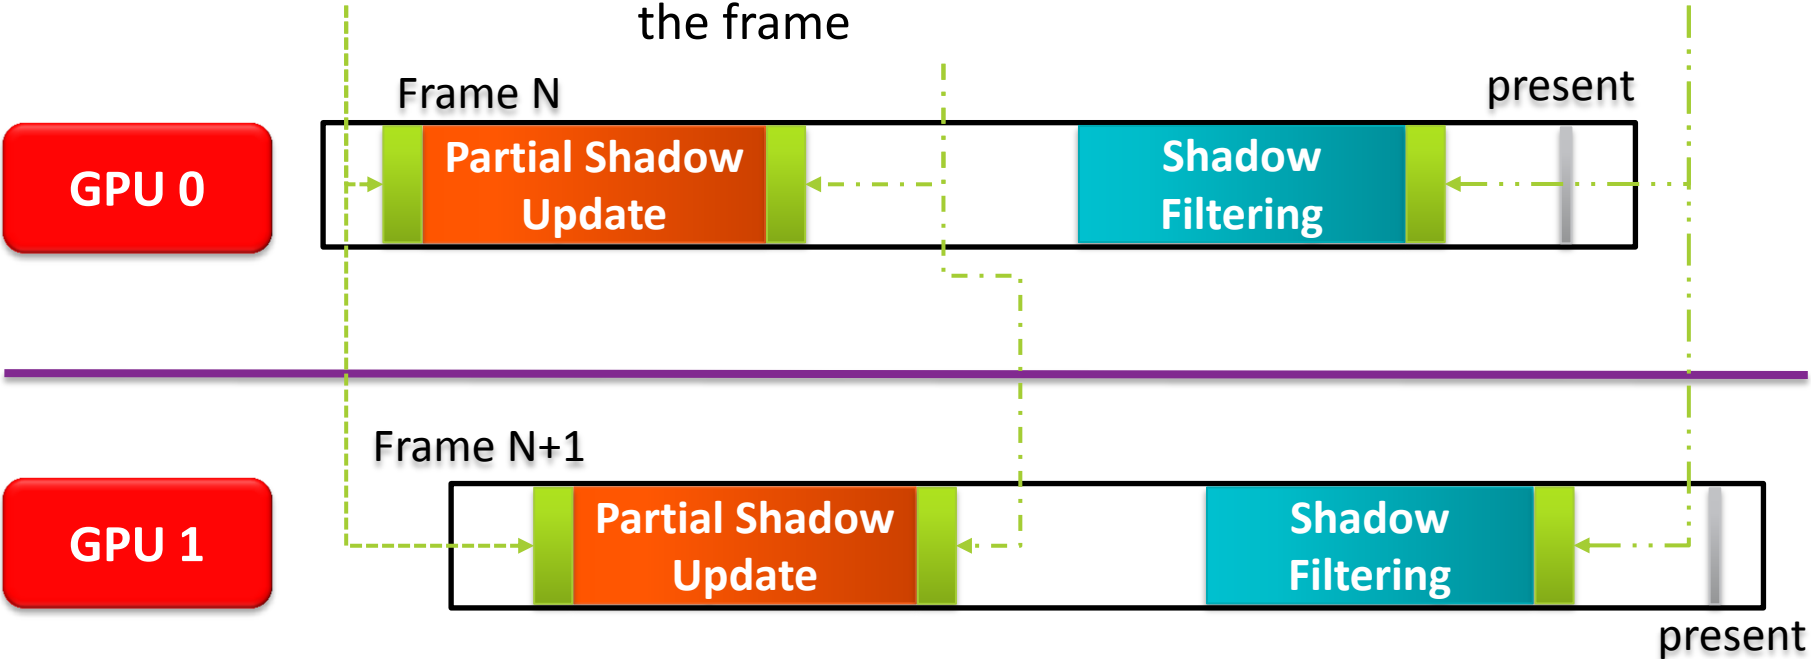
IF RESOURCE IS PARTIALLY UPDATED

- Signal **BeginAllAccess()**

Before the first use of a resource (read or write)
- Signal **EndWrite()**

After the resource has been modified for the last time in the frame
- Signal **EndAllAccess()**

After the resource has been used for the last time within a frame



Why **3** API calls for transfer API?..

AMD CROSSFIRE API

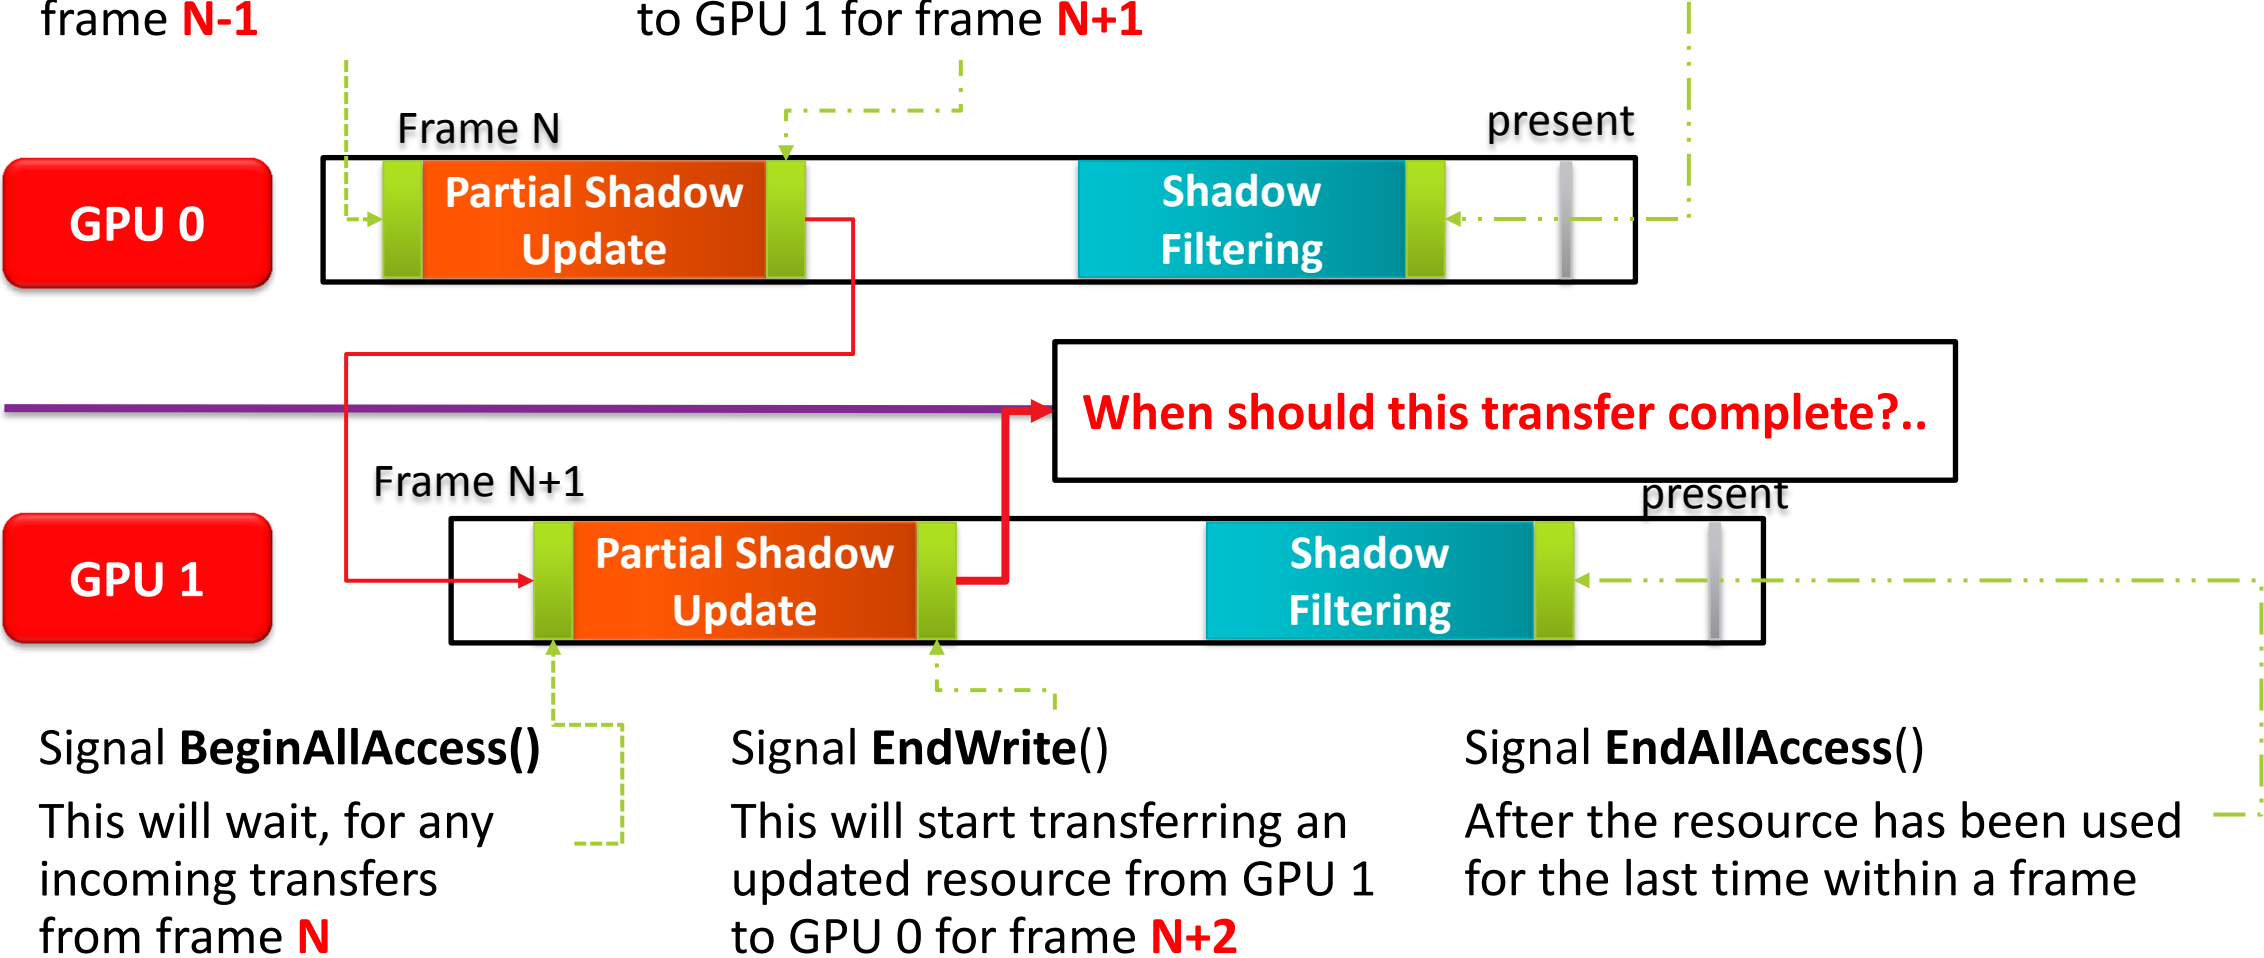


IF RESOURCE IS PARTIALLY UPDATED

Signal **BeginAllAccess()**
This will wait, for any incoming transfers from frame **N-1**

Signal **EndWrite()**
This will start transferring an updated resource from GPU 0 to GPU 1 for frame **N+1**

Signal **EndAllAccess()**
after the resource has been used for the last time within a frame



AMD CROSSFIRE API



IF RESOURCE IS PARTIALLY UPDATED

Signal **BeginAllAccess()**

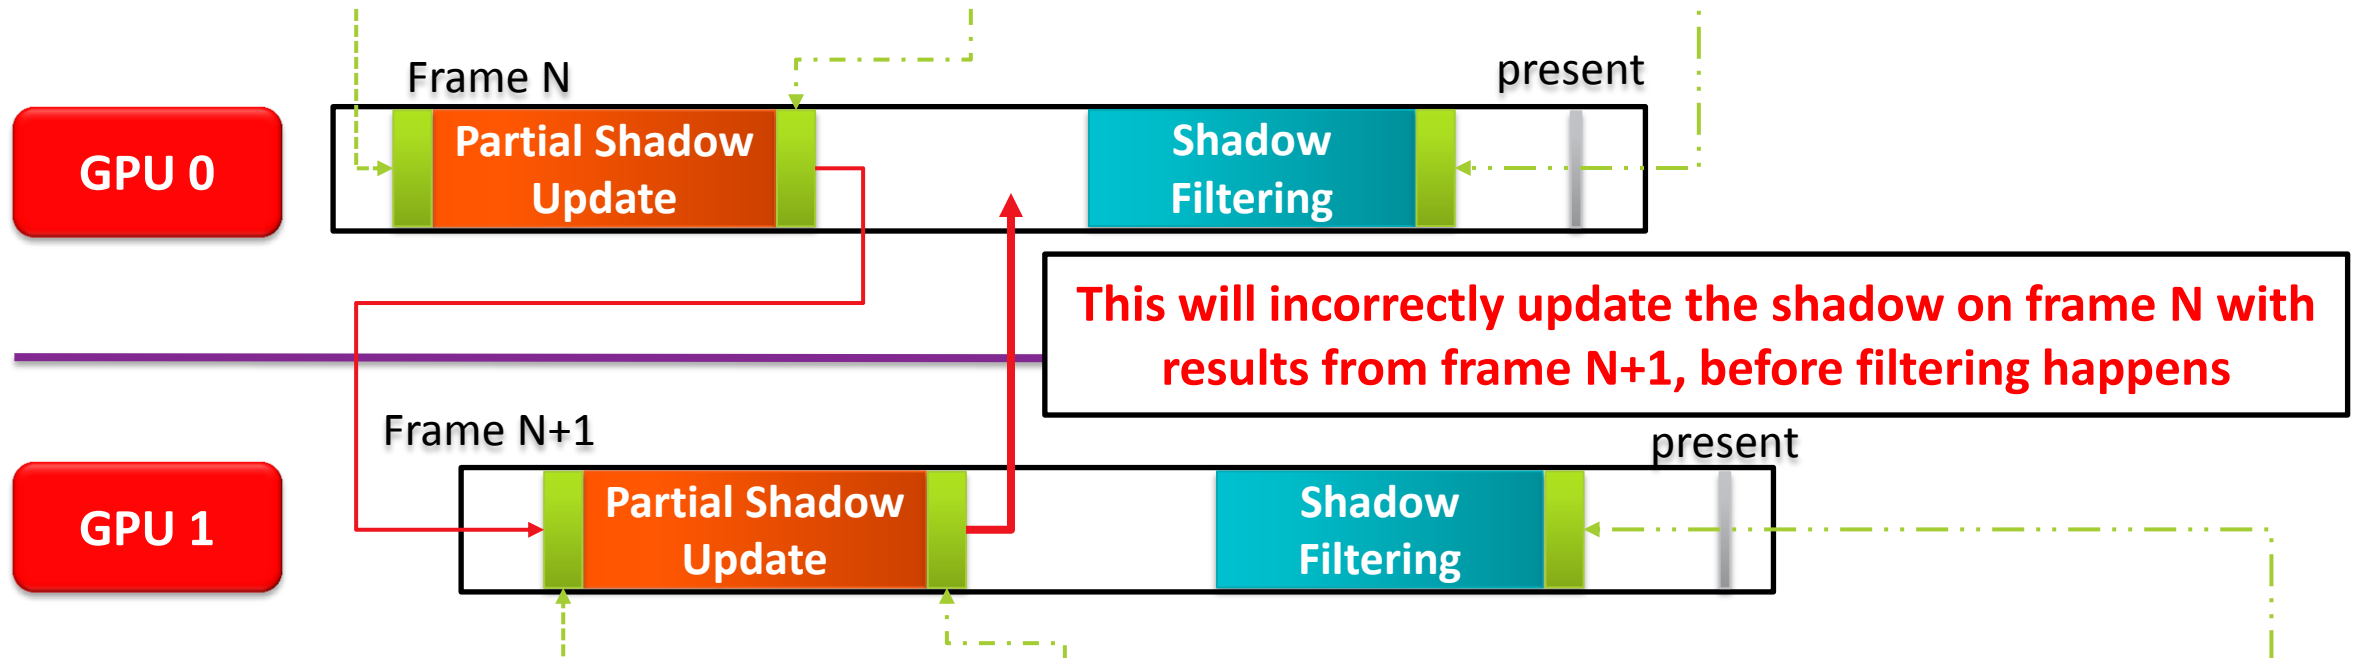
This will wait, for any incoming transfers from frame **N-1**

Signal **EndWrite()**

This will start transferring an updated resource from GPU 0 to GPU 1 for frame **N+1**

Signal **EndAllAccess()**

After the resource has been used for the last time within a frame



Signal **BeginAllAccess()**

This will wait, for any incoming transfers from frame **N**

Signal **EndWrite()**

This will start transferring an updated resource from GPU 1 to GPU 0 for frame **N+2**

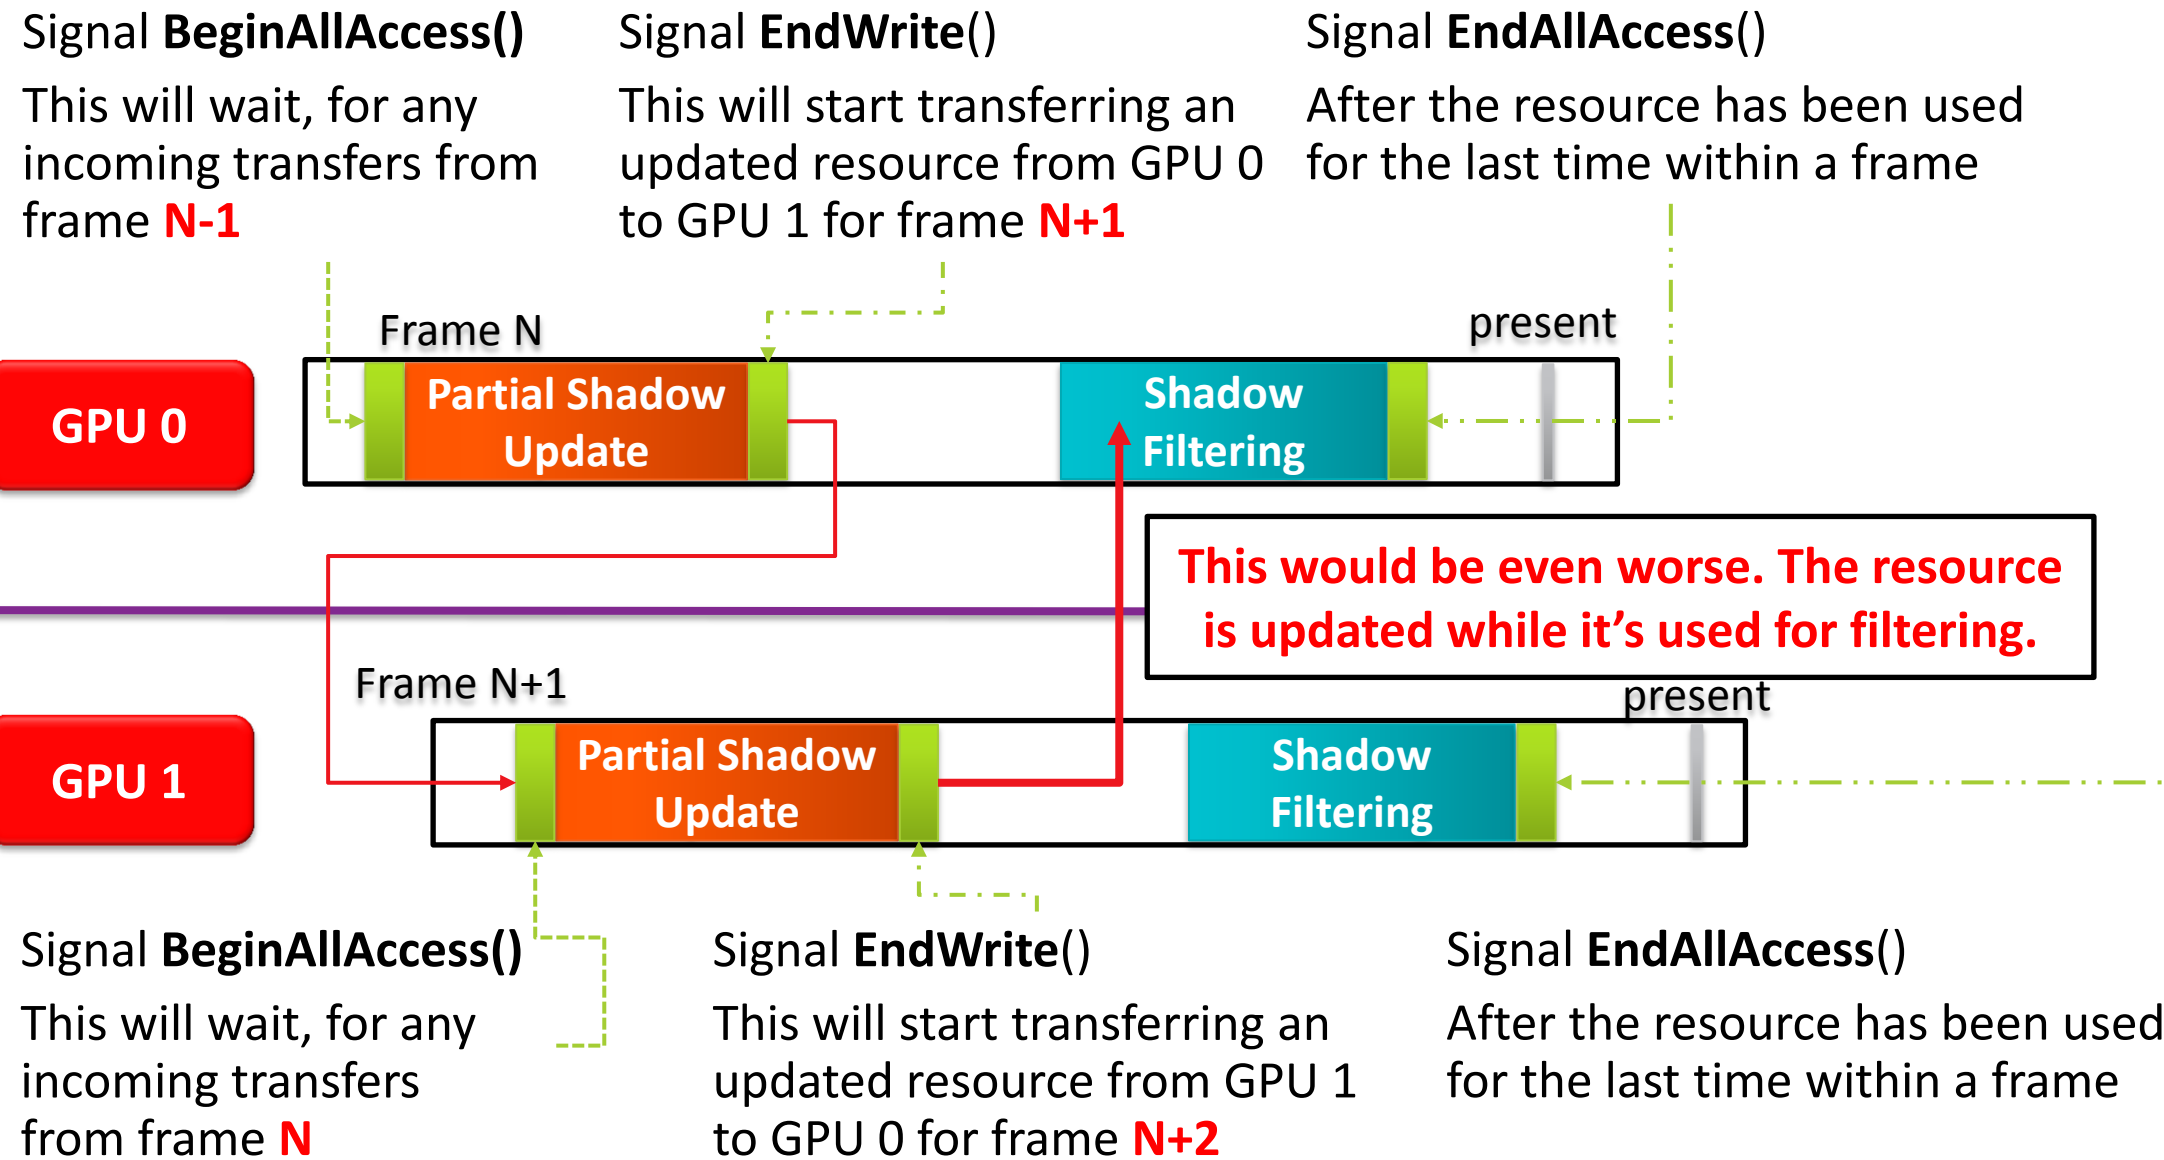
Signal **EndAllAccess()**

After the resource has been used for the last time within a frame

AMD CROSSFIRE API



IF RESOURCE IS PARTIALLY UPDATED



AMD CROSSFIRE API



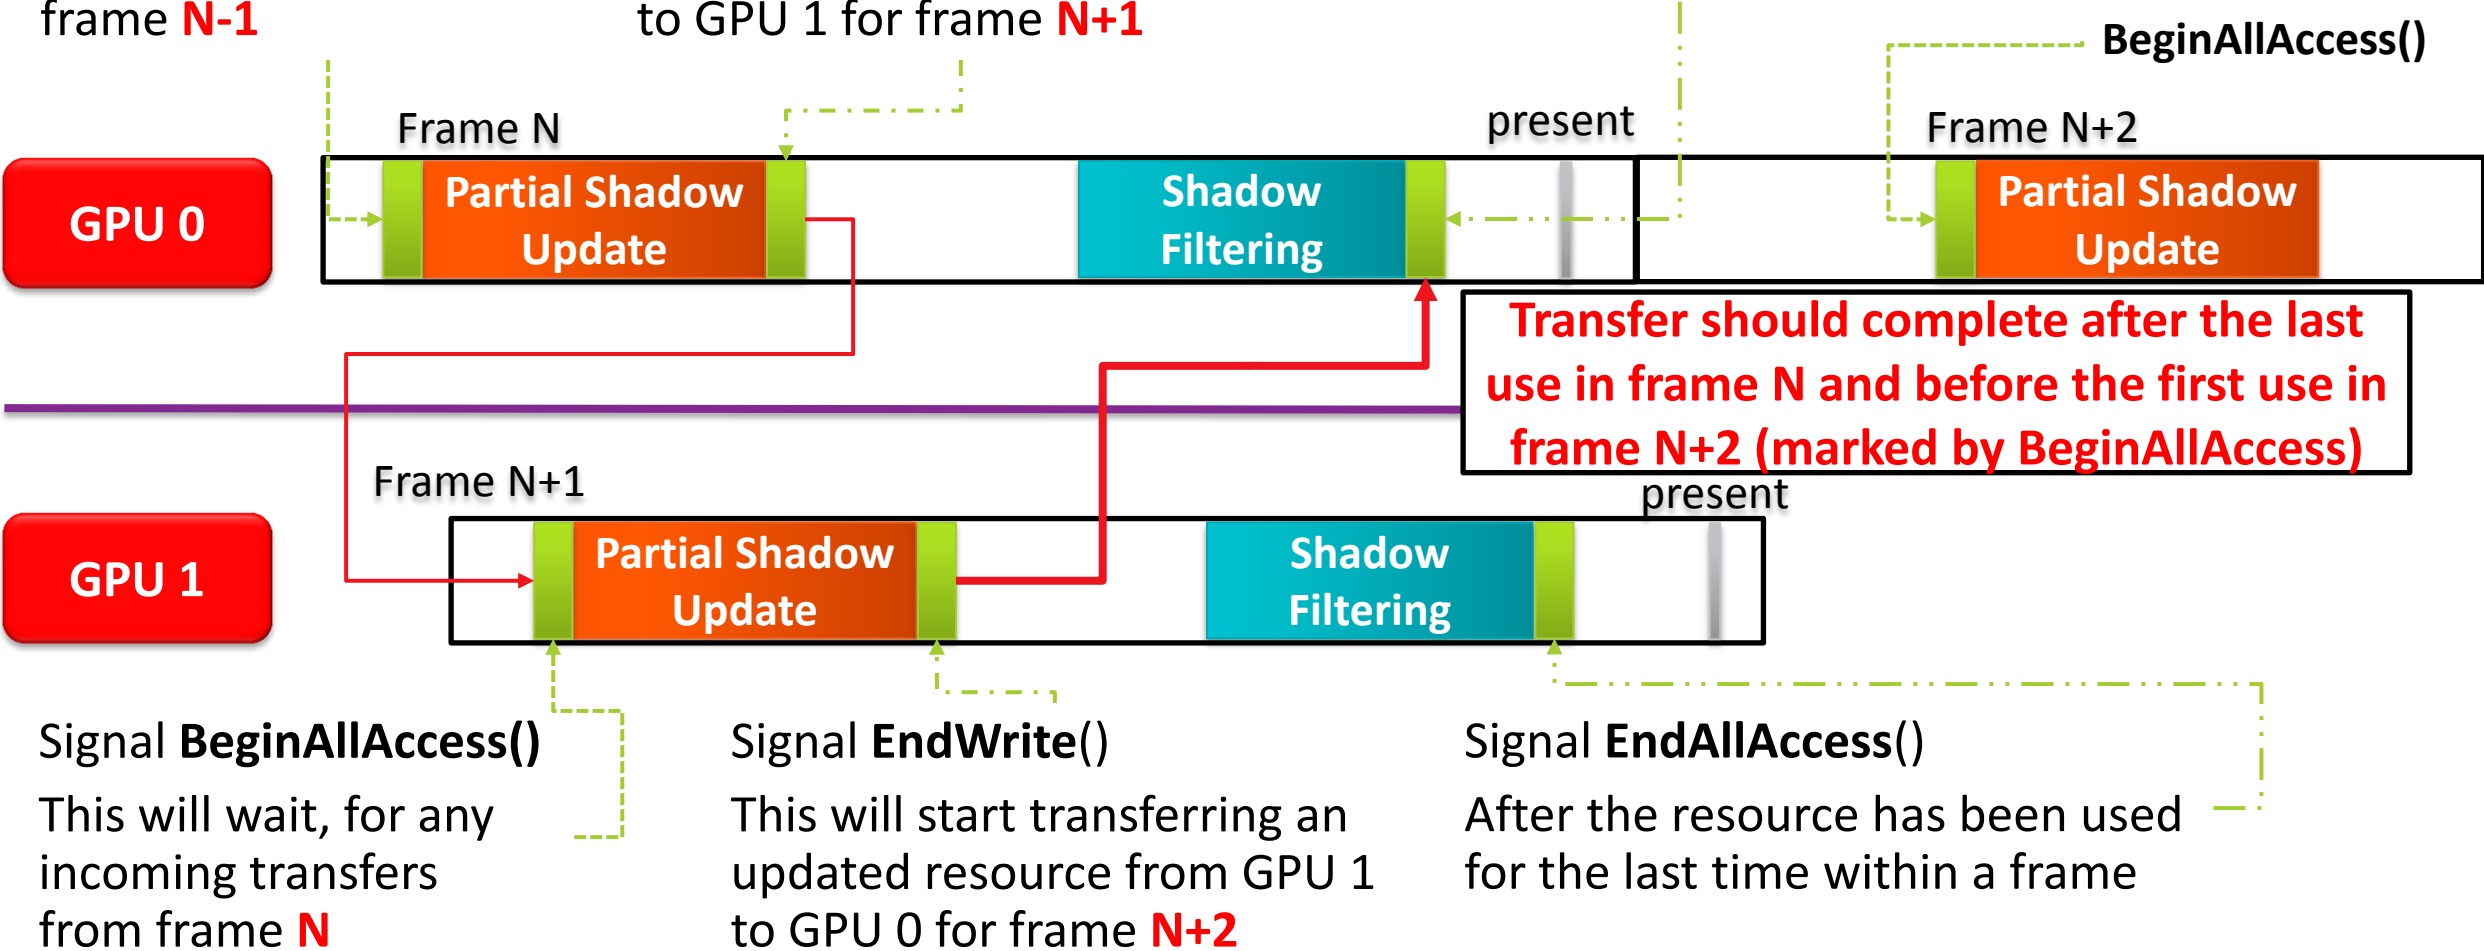
IF RESOURCE IS PARTIALLY UPDATED

- Signal **BeginAllAccess()**

This will wait, for any incoming transfers from frame **N-1**
- Signal **EndWrite()**

This will start transferring an updated resource from GPU 0 to GPU 1 for frame **N+1**
- Signal **EndAllAccess()**

After the resource has been used for the last time within a frame



- ▲ EndWrite – starts a transfer
 - This is optional, if a resource isn't updated within a frame don't initiate transfers
- ▲ BeginAllAccess – waits on any previous transfers to complete
- ▲ EndAllAccess – gates any incoming transfers until resource can be safely updated
- ▲ Time to perform a resource transfer is from EndAllAccess until the next BeginAllAccess
 - \geq (from Present until next BeginAllAccess)
 - EndAllAccess can be naively called at Present()
 - If EndAllAccess is placed as early as the engine allows this maximizes available time for transfer

AMD CROSSFIRE API

TRANSFER_2STEP_GPU



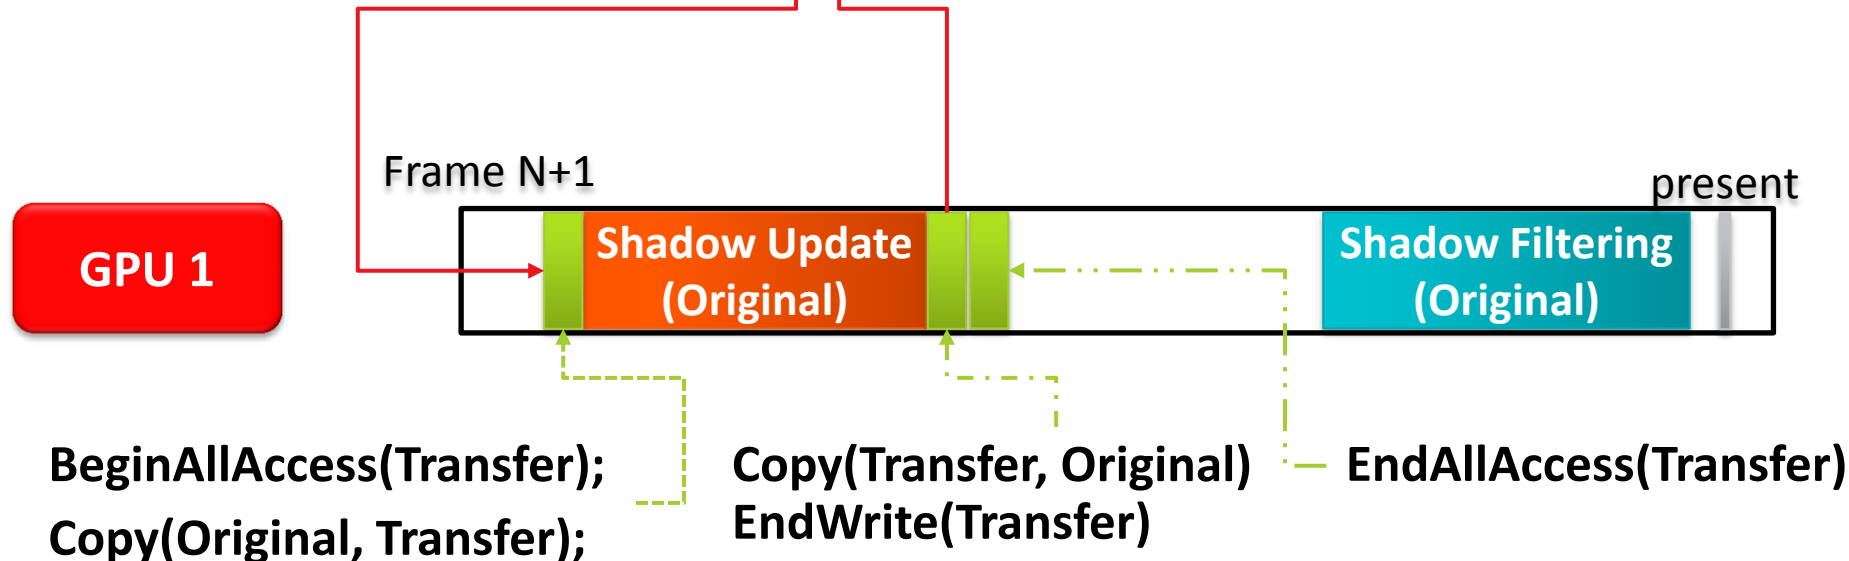
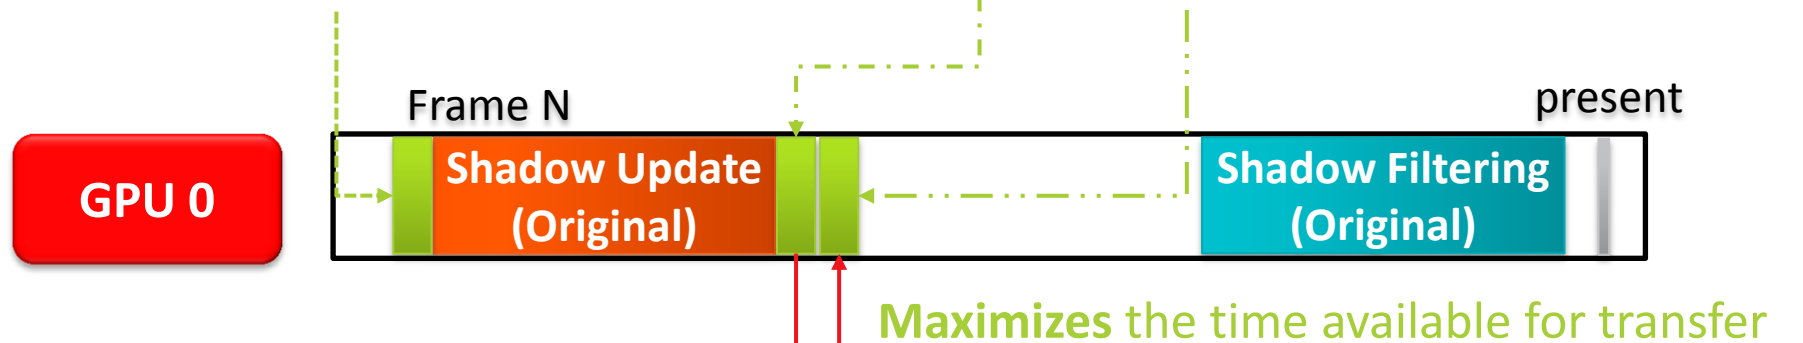
- ▲ Transfer API has a clear dependency on how the resource is used
 - 2 STEP GPU TRANSFER is an optimization that allows to remove dependencies
 - Also simplifies CFX API integration
- ▲ General idea:
 - Given a resource that is updated irregularly
 - it is referred to as the “Original” resource
 - Create “Original” resource with DISABLE transfer flag
 - Create a copy of that resource (same format and size) with applicable transfer flag (1STEP, 2STEP or 2STEP_BROADCAST)
 - It is referred to as the “Transferred” resource
 - On frame *N* when “Original” needs to be updated:
 - BeginAllAccess(“Transferred”) → if there were any updates from previous frames, wait for them
 - Copy(“Original”, Transferred) → if there were any updates from previous frames, they need to be copied into “Original” now
 - Update(“Original”)
 - Copy(“Transferred”, “Original”) → whatever was updated in “Original” needs to be copied into “Transferred”
 - EndWrites(“Transferred”) → “Transferred” is the resource we pass around
 - EndAllAccess(“Transferred”) → Because we know “Transferred” will not be used throughout the frame we are done with it!
- ▲ This is not an API flag, it’s implemented directly in SW and uses an applicable API flag

AMD CROSSFIRE API



TRANSFER_2STEP_GPU

BeginAllAccess(Transfer);
Copy(Original, Transfer);
Copy(Transfer, Original);
EndWrite(Transfer)
EndAllAccess(Transfer)






CROSSFIRE DIRECTX® 11 SAMPLE



CROSSFIRE DIRECTX® 11 SAMPLE

UI



The sample renders a cube shadow map for a point light. This control allows to select the layout of cube faces in memory: a texture2d atlas or a texture2d array

Each face of a cube map is of "Shadow Map Size" resolution which can be modified via this control

This check box enables the application to use the crossfire API. Otherwise applications behaviour is completely determined by the Radeon Settings app profile

This check box enables an optimization that is called "TRANSFER_2STEP_GPU". Similar to the CFX API flags TRANSFER_2STEP_* a "TRANSFER_2STEP_GPU" first copies the target resource into an auxiliary GPU resource, which is later used for the actual transfer. In contrast to API flags this optimization is implemented directly in the sample, using the API.

This check box enables the application to delay calling EndAllAccess() until the end of the frame. This can be used to measure performance advantage of transferring a resource in advance, in contrast to transferring the resource at the end of the frame.

This check box allows to control sample's strategy for shadow map updates. The sample renders a cube shadow map and can either:

- render all 6 shadow map faces once in 6 frames and initiate a broadcast of the full shadow map resource only once
- update a single face of cube shadow map each frame and broadcast just that subregion (for atlas) or subresource (for array)




CROSSFIRE DIRECTX® 11 SAMPLE




UI

- ▲ Sample renders a cube shadow map for a point light
 - A total of 6 shadow faces
- ▲ “Shadow Map Size” controls the size of a single cube face
 - So 1024^2 really means 6×1024^2
- ▲ Shadow map layout
 - A 2D texture atlas
 - A 2D texture array, each cube face located in a separate array slice
- ▲ Shadow map updating schemes
 - Update all 6 shadow faces once in 6 frames
 - Frame I where $(I \% 6) == 0$ renders more geometry and is slower
 - Frame I where $(I \% 6) == 0$ broadcasts the full resource
 - Update 1 shadow face each frame
 - Each frame broadcasts a single sub-region (or a sub-resource) of the full resource
 - On average, both updating schemes:
 - render the same amount of geometry
 - transfer the same amount of data

Texture2D Atlas Layout



Texture2D Array Layout





PERFORMANCE 



PERFORMANCE CHARACTERISTICS (2 RADEON R9 290X)



FOREWORD

- ▲ Numeric values in the performance tables are FPS
- ▲ Test Hardware
 - Two Radeon R9 290X
 - Core i5 3330
- ▲ Expectations for default driver modes (when CFX API is not used):
 - AFR-Friendly mode
 - **Flicker** is **expected** for both shadow map update schemes
 - This profile provides us with the Speed of Light (SOL) Crossfire performance
 - AFR-Compatible mode
 - **Flicker** is **expected** when 1 shadow face is updated each frame
 - The driver fails to detect a stale resource. This is explained on **slides 6-7**
 - This profile provides us with default Crossfire transfer performance
 - Crossfire API transfer performance is expected to be higher because sample initiates transfers earlier
- ▲ Crossfire API was tested with Radeon Settings Crossfire profile for the sample set to AFR-Friendly
 - Crossfire API performance is expected to be : AFR-Friendly \geq Crossfire API \geq^* AFR-Compatible
 - 2-STEP-GPU-TRANSFER has overhead for resource copies
 - * AFR-Compatible mode **will be as fast as AFR-Friendly** for shadow map update scheme {1 shadow face updated each frame) because the driver doesn't detect a stale resource.

PERFORMANCE CHARACTERISTICS



TRANSFER FLAG: 2STEP_WITH_BROADCAST

Texture2D Atlas (each subregion 512 ²)	Single GPU	AFR Friendly	AFR Compatible	Crossfire API	Crossfire API: 2-STEP-GPU TRANSFER	Crossfire API: Delay EndAllAccess	Crossfire API 2-STEP- GPU TRANSFER + Delay EndAllAccess
6 shadow maps / updates once every 6 frames	274	352 Flicker	294	298	318	294	298
1 shadow map updated every frame	273	468 Flicker	466 Flicker	235	428	235	419
Texture2D Array (subresource 512 ²)	Single GPU	AFR Friendly	AFR Compatible	Crossfire API	Crossfire API: 2-STEP-GPU TRANSFER	Crossfire API: Delay EndAllAccess	Crossfire API 2-STEP- GPU TRANSFER + Delay EndAllAccess
6 shadow maps / updates once every 6 frames	273	351 Flicker	297	294	300	289	276
1 shadow map updated every frame	273	465 Flicker	455 Flicker	405	400	402	390

PERFORMANCE CHARACTERISTICS



TRANSFER FLAG: 2STEP_WITH_BROADCAST

Texture2D Atlas (each subregion 1024 ²)	Single GPU	AFR Friendly	AFR Compatible	Crossfire API	Crossfire API: 2-STEP-GPU TRANSFER	Crossfire API: Delay EndAllAccess	Crossfire API 2-STEP- GPU TRANSFER + Delay EndAllAccess
6 shadow maps / updates once every 6 frames	270	346 Flicker	130	178	209	179	209
1 shadow map updated every frame	268	449 Flicker	450 Flicker	95	342	95	336
Texture2D Array (subresource 1024 ²)	Single GPU	AFR Friendly	AFR Compatible	Crossfire API	Crossfire API: 2-STEP-GPU TRANSFER	Crossfire API: Delay EndAllAccess	Crossfire API 2-STEP- GPU TRANSFER + Delay EndAllAccess
6 shadow maps / updates once every 6 frames	269	347 Flicker	130	174	169	173	169
1 shadow map updated every frame	269	449 Flicker	448 Flicker	303	297	307	297

PERFORMANCE CHARACTERISTICS



TRANSFER FLAG: 2STEP_WITH_BROADCAST

Texture2D Atlas (each subregion 2048 ²)	Single GPU	AFR Friendly	AFR Compatible	Crossfire API	Crossfire API: 2-STEP-GPU TRANSFER	Crossfire API: Delay EndAllAccess	Crossfire API 2-STEP- GPU TRANSFER + Delay EndAllAccess
6 shadow maps / updates once every 6 frames	256	322 Flicker	40	74	71	73	73
1 shadow map updated every frame	251	424 Flicker	424 Flicker	25	123	25	123
Texture2D Array (subresource 2048 ²)	Single GPU	AFR Friendly	AFR Compatible	Crossfire API	Crossfire API: 2-STEP-GPU TRANSFER	Crossfire API: Delay EndAllAccess	Crossfire API 2-STEP- GPU TRANSFER + Delay EndAllAccess
6 shadow maps / updates once every 6 frames	256	320 Flicker	40	81	74	80	74
1 shadow map updated every frame	257	429 Flicker	425 Flicker	136	121	135	120

PERFORMANCE CHARACTERISTICS



TRANSFER FLAG: 2STEP_WITH_BROADCAST

Texture2D Atlas (each subregion 4096 ²)	Single GPU	AFR Friendly	AFR Compatible	Crossfire API	Crossfire API: 2-STEP-GPU TRANSFER	Crossfire API: Delay EndAllAccess	Crossfire API 2-STEP- GPU TRANSFER + Delay EndAllAccess
6 shadow maps / updates once every 6 frames	218	254 Flicker	12	22	20	22	20
1 shadow map updated every frame	198	348 Flicker	250 Flicker	7	32	7	32
Texture2D Array (subresource 4096 ²)	Single GPU	AFR Friendly	AFR Compatible	Crossfire API	Crossfire API: 2-STEP-GPU TRANSFER	Crossfire API: Delay EndAllAccess	Crossfire API 2-STEP- GPU TRANSFER + Delay EndAllAccess
6 shadow maps / updates once every 6 frames	218	256 Flicker	12	25	23	25	23
1 shadow map updated every frame	218	355 Flicker	361 Flicker	36	32	36	32

DISCLAIMER & ATTRIBUTION



The information presented in this document is for informational purposes only and may contain technical inaccuracies, omissions and typographical errors.

The information contained herein is subject to change and may be rendered inaccurate for many reasons, including but not limited to product and roadmap changes, component and motherboard version changes, new model and/or product releases, product differences between differing manufacturers, software changes, BIOS flashes, firmware upgrades, or the like. AMD assumes no obligation to update or otherwise correct or revise this information. However, AMD reserves the right to revise this information and to make changes from time to time to the content hereof without obligation of AMD to notify any person of such revisions or changes.

AMD MAKES NO REPRESENTATIONS OR WARRANTIES WITH RESPECT TO THE CONTENTS HEREOF AND ASSUMES NO RESPONSIBILITY FOR ANY INACCURACIES, ERRORS OR OMISSIONS THAT MAY APPEAR IN THIS INFORMATION.

AMD SPECIFICALLY DISCLAIMS ANY IMPLIED WARRANTIES OF MERCHANTABILITY OR FITNESS FOR ANY PARTICULAR PURPOSE. IN NO EVENT WILL AMD BE LIABLE TO ANY PERSON FOR ANY DIRECT, INDIRECT, SPECIAL OR OTHER CONSEQUENTIAL DAMAGES ARISING FROM THE USE OF ANY INFORMATION CONTAINED HEREIN, EVEN IF AMD IS EXPRESSLY ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

ATTRIBUTION

© 2013 Advanced Micro Devices, Inc. All rights reserved. AMD, the AMD Arrow logo and combinations thereof are trademarks of Advanced Micro Devices, Inc. in the United States and/or other jurisdictions. Other names are for informational purposes only and may be trademarks of their respective owners.